

SHUVAM PAL

+91 85830 54679 ✉ ishuvam.pal@gmail.com  Shuvam Pal  unichronic

Education

Dayananda Sagar College of Engineering

Bengaluru, India

B.E. Artificial Intelligence and Machine Learning

Experience

Bentham AI 

Dec 2025 – Feb 2026

Backend Engineering Intern: Go, Node.js, Docker, GCP, Puppeteer, GitHub Actions

Remote

- Built a resilient **Puppeteer/Node.js workflow** for MCA filings with session handling, CAPTCHA handoff, retries, and state recovery, reducing manual filing time by **85%**.
- Optimized an **AI-assisted company-registration pipeline** for unique company-name suggestions, descriptions, and filing-field drafts, reducing manual drafting/rework across MCA workflow preparation.
- Containerized backend services with Docker, wired GitHub Actions deployments to Google Cloud Run, and improved release velocity by **70%**.

Google Summer of Code 

May 2025 – Sept 2025

Mentee at Invesalious: Python, VTK, wxPython, NumPy, PyTorch, ONNX

Remote

- Built Python medical-image segmentation tooling for **95 brain subparts**, handling preprocessing, orientation correction, model inference, generated masks, and large-volume data paths.
- Reduced segmentation runtime by **87%** using asynchronous and parallel execution while preserving inspection workflows for generated masks and 3D surfaces.
- Debugged complex pipeline issues across conformation, orientation, label mapping, PyTorch/ONNX-style inference, and binary mask generation in a mature open-source codebase.

Open Source

[Contributions](#) 

- Contributed patches across Invesalious, IOOS, VideoLAN, CRIU, Tiled, Kyverno, and Kuadrant, working inside maintainer review processes for reliability, security, and integration quality.

Projects

Swish Support System | *Python, Go, FastAPI, Gin, PostgreSQL, Redis, Langfuse*

[GitHub](#) 

- Built an **AI-assisted support backend** where LLM calls classify messy customer complaints, intent, severity, ambiguity, and photo-evidence usefulness while deterministic code controls refund, replacement, coupon, and escalation policy.
- Maintained per-conversation context, evidence artifacts, order/customer state, and traceable actions across multi-turn workflows with PostgreSQL, Redis, Langfuse/local traces, and regression tests for edge cases.

Murdoc | *Python, FastAPI, MCP, OpenTelemetry, Prometheus, Docker*

[GitHub](#) 

- Instrumented an AI gateway with OpenTelemetry traces, Prometheus metrics, route profiles, health checks, and audit records tying model/tool requests to policy outcomes.
- Built replayable decision records for prompt, tool, retrieval, and output layers so high-risk AI behavior could be correlated with request context and failure modes.

Seaweed Contest Platform | *Go, PostgreSQL, Redis, Judge0, AWS, Docker, Kubernetes*

[GitHub](#) 

- Built backend services for registration, code submissions, sandboxed Judge0 execution, scoring, and live contest rankings across **500+ concurrent users**.
- Designed async Go dispatch, verdict polling, PostgreSQL result schemas, Redis-backed runtime state, and Docker/Kubernetes deployment on AWS, reaching **sub-5s P95 verdict latency**.

Technical Skills

Languages: Python, Go, TypeScript, C++

Frameworks: FastAPI, Gin, Echo, Node.js, Agno, PyTorch

Tools: PostgreSQL, Redis, Docker, Kubernetes, AWS, GCP, OpenTelemetry, Prometheus, Langfuse, GitHub Actions